

# Getting started with CUDA

## Part 1 - CUDA overview

---

Edwin Carlinet, Joseph Chazalon

`firstname.lastname@epita.fr`

Fall 2023

EPITA Research Laboratory (LRE)





# CUDA overview

---

# What is CUDA?

A product

- It enables to use NVidia GPUs for computation

A C/C++ variant

- Mostly C++14-compatible, with extensions
- and also some restrictions!

A SDK

- A set of compilers and toolchains for various architectures
- Performance analysis tools

A runtime

- An assembly specification
- Computation libraries (linear algebra, etc.)

A new industry standard

- Used by every major deep learning framework
- Replacing OpenCL as Vulkan is replacing OpenGL

# The CUDA ecosystem (missing L and H series)

## GPU Computing Applications

### Libraries and Middleware

cuDNN TensorRT	cuFFT cuBLAS cuRAND cuSPARSE	CUDA MAGMA	Thrust NPP	VSIP SVM OpenCurrent	PhysX OptiX iRay	MATLAB Mathematica
-------------------	---------------------------------------	---------------	---------------	----------------------------	------------------------	-----------------------

### Programming Languages

C	C++	Fortran	Java Python Wrappers	DirectCompute	Directives (e.g. OpenACC)
---	-----	---------	----------------------------	---------------	------------------------------



### CUDA-Enabled NVIDIA GPUs

NVIDIA Ampere Architecture (compute capabilities 8.x)				Tesla A Series
NVIDIA Turing Architecture (compute capabilities 7.x)		GeForce 2000 Series	Quadro RTX Series	Tesla T Series
NVIDIA Volta Architecture (compute capabilities 7.x)	DRIVE/JETSON AGX Xavier		Quadro GV Series	Tesla V Series
NVIDIA Pascal Architecture (compute capabilities 6.x)	Tegra X2	GeForce 1000 Series	Quadro P Series	Tesla P Series



# The CUDA ecosystem (not so long ago)

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIP, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier	GeForce 2000 Series	Quadro RTX Series	Tesla T Series		
Volta Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier			Tesla V Series		
Pascal Architecture (Compute capabilities 6.x)	Tegra X2	GeForce 1000 Series	Quadro P Series	Tesla P Series		
Maxwell Architecture (Compute capabilities 5.x)	Tegra X1	GeForce 900 Series	Quadro M Series	Tesla M Series		
Kepler Architecture (Compute capabilities 3.x)	Tegra K1	GeForce 700 Series GeForce 600 Series	Quadro K Series	Tesla K Series		
	EMBEDDED	CONSUMER DESKTOP, LAPTOP	PROFESSIONAL WORKSTATION	DATA CENTER		

Figure 2: The CUDA ecosystem

# Libraries *and* Compiler Directives *and* Programming Language

CUDA is mostly based on a “new” **programming language**: CUDA C (or C++, or Fortran).  
*This grants much flexibility and performance*

But it also exposes much of GPU goodness through **libraries**.

And it supports a few **compiler directives** to facilitate some constructs.

```
#pragma unroll  
for(int i = 0; i < WORK_PER_THREAD; ++i)  
    // Some thread work
```

# The big idea: Kernels instead of loops


## Without CUDA (vector addition)

```
// compute vector sum C = A + B
void vecAdd(float *h_A, float *h_B, float *h_C, int n)
{
    for (int i = 0; i < n; ++i)
        h_C[i] = h_A[i] + h_B[i];
}

int main()
{
    // MISSING: Allocation for A, B and C
    // MISSING: I/O to read n elements of A and B
    vecAdd(h_A, h_B, h_C, n);
}
```



## With CUDA (1/2): move work to the separate compute device



**Figure 3:** Computation on separate device

```
#include <cuda.h>
void vecAdd(float *h_A, float *h_B, float *h_C, int n)
{
    int size_bytes = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    // 1.1 Allocate device memory for A, B and C
    // 1.2 Copy A and B to device memory

    // 2. Launch kernel code - computation done on device

    // 3. Copy C (result) from device memory
    // Free device vectors
}

int main() { /* Unchanged */ }
```

## With CUDA (2/2): Kernel sample code

```
// kernel  
__global__ void kvecAdd(float *d_A, float *d_B, float *d_C, int n)  
{  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i >= n) { return; }  
    d_C[i] = d_A[i] + d_B[i];  
}
```

**No more for loop!**

# Arrays of parallel threads

A CUDA kernel is executed by a **grid** (array) of threads

- All threads in a grid run the same kernel code (Single Program Multiple Data)
- Each thread has indexes that it uses to compute memory addresses and make control decisions




Figure 4: A thread block

# Thread blocks

Threads are grouped into thread blocks

- Threads within a block cooperate via
  - **shared memory**
  - **atomic operations**
  - **barrier synchronization**
- Threads in different blocks do not interact<sup>1</sup>

Thread block 1    Thread block 2    ...    Thread block N-1




Figure 5: Independent thread blocks

<sup>1</sup>Not in this course, though there are techniques for that.

## A multidimensional grid of computation threads (1/2)

Each thread uses indices (added by the compiler) to decide what data to work on:

- `blockIdx` ( $0 \rightarrow \text{gridDim}$ ): 1D, 2D or 3D
- `threadIdx` ( $0 \rightarrow \text{blockDim}$ ): 1D, 2D or 3D

Each index has `x`, `y` and `z` attributes to get the actual index in each dimension.


```
int i = threadIdx.x;  
int j = threadIdx.y;  
int k = threadIdx.z;
```

Simplifies memory addressing when processing multidimensional data:

- *image processing*
- *solving PDE on volumes*
- ...

## A multidimensional grid of computation threads (2/2)

Grid and blocks can have different dimensions, but they usually are two levels of the same work decomposition.



**Figure 6:** An example of 2D grid with 3D blocks

## Grid & block examples (1/2)

### Vector addition (N elements)

*// Kernel definition*

```
__global__ void VecAdd(float* d_A, float* d_B, float* d_C, int sz)
{
    int i = threadIdx.x;  // /\ Assuming 1 block here
    if (i >= sz) { return ; }
    d_C[i] = d_A[i] + d_B[i];
}
```

```
int main()
```

```
{
    ...
    // Kernel invocation with N threads in a single block
    VecAdd<<<1, N>>>(A, B, C, sz);  // <-- So this is how we launch CUDA kernels!
    ...
}
```

## Grid & block examples (2/2)

### Matrix addition ( $N \times N$ elements)

*// Kernel definition*

```
__global__ void MatAdd(float d_A[N][N], float d_B[N][N], float d_C[N][N], int sz)
{
    int i = threadIdx.x; // /\ Assuming 1 block here
    int j = threadIdx.y; // /\ Assuming 1 block here
    if (i >= sz || j >= sz) { return; }
    d_C[i][j] = d_A[i][j] + d_B[i][j];
}
```

```
int main()
{
    ...
    int numBlocks = 1; // grid size: 1 * 1 * 1 blocks
    dim3 threadsPerBlock(N, N); // block size: N * N * 1 threads
    MatAdd<<<numBlocks, threadsPerBlock>>>(A, B, C, sz);
    ...
}
```



# Block decomposition enable automatic scalability

Because the work is divided into **independent blocs** which can be **run in parallel** on each *streaming multiprocessor* (SM), the same code can be **automatically scaled** to architectures with more or less SMs. . .

as long as SMs architectures are compatibles (100% compatible with the same Compute Capabilities version — a family of devices, careful otherwise).

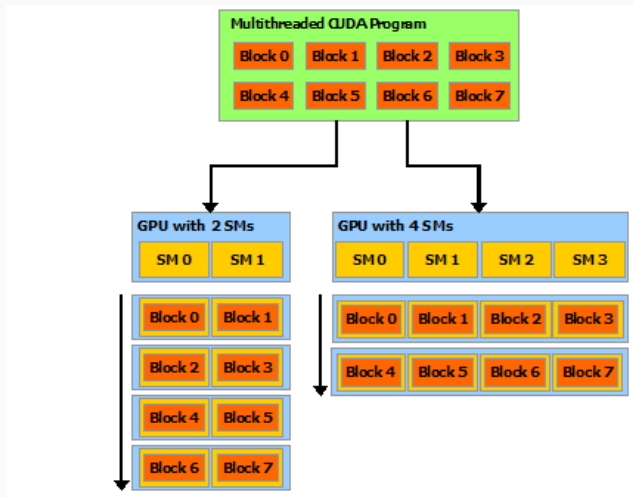


Figure 7: Automatic scaling

# Building and running a simple program

## CUDA Hello world (hello.cu)

```
#include <stdio.h>

__global__ void print_kernel() {
    printf(
        "Hello from block %d, thread %d\n",
        blockIdx.x, threadIdx.x);
}

int main() {
    print_kernel<<<2, 3>>>();
    cudaDeviceSynchronize();
}
```

## Compile

```
$ nvcc hello.cu -o hello
```

## Run

```
$ ./hello
Hello from block 1, thread 0
Hello from block 1, thread 1
Hello from block 1, thread 2
Hello from block 0, thread 0
Hello from block 0, thread 1
Hello from block 0, thread 2
```

# What you need to get started

NVidia GPU hardware

NVidia GPU drivers, properly loaded

*modprobe nvidia ...*

CUDA runtime libraries

*libcuda.so, libnvidia-fatbinaryloader.so, ...*

CUDA SDK (NVCC compiler in particular)

*relies on a standard C/C++ compiler and toolchain*

*[docs.nvidia.com/cuda/cuda-installation-guide-linux](https://docs.nvidia.com/cuda/cuda-installation-guide-linux)*

Basic C/C++ knowledge

## Host vs Device $\leftrightarrow$ Separate memory

*GPUs are computation units which require explicit usage, as opposed to a CPU*

*Need to load data to and fetch result from device*

## Replace loops with kernels

*Kernel = Function computed in relative isolation on small chunks of data, on the GPU*

## Divide the work

*Problem  $\rightarrow$  Grid  $\rightarrow$  Blocks  $\rightarrow$  Threads*

## Compile and run using CUDA SDK

*nvcc, libcuda.so, ...*